Assignment 1 Report

Julius Putra Tanu Setiaji (A0149787E), Chen Shaowei (A0110560Y)

1 October 2018

1 Program Design

This train simulation is implemented in OpenMP. Primary design considerations are:

- Each train is simulated by one thread as required in the specifications.

- This is achieved by calling `omp_set_num_threads(num_trains);`

- Each train stores its next state, and when it will enter that state.

- Each train will generate the time at which it will open its doors when it starts waiting at a station.

Assumptions

- Only one train can open its doors at each at any one time, regardless of direction.

- Train stations have infinite capacity for waiting trains.

- Time units are discrete and can have no subdivisions

  - **Implication**: It is sufficient to store all time units as integers instead of floating point numbers

- Trains must open their doors for at least 1 unit of time.

  - **Implication**: We round every randomly generated door open time up to the nearest integer

2 Points to Note / Implementation Details

- Current simulation time needs to be shared across all threads. In addition, time can only be advanced after all threads have completed the actions to be done in the current tick.

- Each train is a finite state machine with 4 states: `OPEN_DOOR`, `CLOSE_DOOR`, `DEPART` and `ARRIVE`.

- Each train keeps track of its next action time (time for it to change to the next state), and actions to be completed within the current tick are performed in a `while` loop.

- There is a `#pragma omp barrier` to ensure that all threads exit the `while` loop before time is advanced.

- The advancement of time is done in a `#pragma omp single` block to ensure that it is only performed by one thread. In addition, `#pragma omp single` has an implicit barrier at the end of the block so this prevents other threads from entering the next iteration of the `while` loop until the advancement of time is complete.

- Certain resources i.e. train tracks, door-opening rights are limited, and only one train may access them at any time.

  1. One way to implement this is to have threads block when waiting to access such resources. However, this would interfere with the way we have chosen to implement time within this simulation. Blocked threads would prevent code after a `#pragma omp barrier` statement from being executed. We therefore decided **not** to go with this implementation.

  2. An alternative way to implement is to have a queue of trains requesting access to such resources. At each tick, each train would check whether it is at the head of the queue, and if so, it will be able to access the resource. We realised that it will also be necessary to store the time at which a train would be able to gain access to the resource, since our simulation time advances in integer increments, but door open time may have a fractional value. Upon further consideration, we realised that this design could be simplified.

  3. The key insight we arrived at is that any train waiting for access to a resource only needs to know the next time said resource will be available. Since this system does not permit a train to give up waiting for a resource, this can be implemented simply with a thread-safe timekeeper object. When a train requests access to a resource, it tells the timekeeper how much time it will occupy the resource for. The timekeeper will then inform the tra in of the time when the train can access the resource, and update its internal next available time. This is the implementation we decided to go with. In another word, we are implementing an implicit queue for a First-Come-First-Serve (FCFS) scheduling.

- The next consideration is ensuring that system statistics are reported correctly. Since we assume that trains cannot open its doors for 0s, only one train can open its doors at each station per tick. There is therefore no potential race condition in the update of statistics.

- Since print statements are not atomic, we wrapped them in a `#pragma omp critical` block to ensure that only one print operation executes at any one time.

3 Execution Time

3.1 Testcase Used

We use a map generated by a Ruby script which can be found in appendix. The map generated will have at least 4 vertices with degree = 1, and all the paths forming the train lines are at least of length 4. However, these values are configurable.

We ran the same map for 100, 1000, and 10000 time-ticks with all possible different combinations of numbers of trains in each line as long as the total number of trains is between 1 and 64 inclusive. This results in around 47,900 testcases for each time-tick size. As such, we will only include the scatter diagram of the results in the report. However, should the need arise, csv files containing the raw data is attached together with the report.

Below you can find a sample input and visualisation of the adjacency matrix and the train lines. The parameters that we used are number of stations = 10, maximum distance between stations = 10.

Do also note that to facilitate more accurate execution time analysis, we disabled per-tick status output for the following tests.

Figure 1: Sample Input

```
1   10
2   Somerset,Tan Kah Kee,Redhill,Sembawang,Riviera,Samudera,Sengkang,Boon Lay,Tampines
    ↪   West,Tampines
3   0 0 0 0 0 0 0 0 6 0
4   0 0 1 0 0 0 1 0 4 1
5   0 1 0 0 6 0 0 0 0 0
6   0 0 0 0 0 0 0 4 0 9
7   0 0 6 0 0 0 0 0 0 0
8   0 0 0 0 0 0 0 0 0 10
9   0 1 0 0 0 0 0 0 0 0
10  0 0 0 4 0 0 0 0 0 0
11  6 4 0 0 0 0 0 0 0 0
12  0 1 0 9 0 10 0 0 0 0
13  0.1,0.2,0.1,0.1,1.0,0.5,0.7,0.8,0.5,0.5
14  Somerset,Tampines West,Tan Kah Kee,Tampines,Sembawang,Boon Lay
15  Samudera,Tampines,Tan Kah Kee,Sengkang
16  Somerset,Tampines West,Tan Kah Kee,Sengkang
17  10000
18  21,22,21
```

Figure 2: Map of the train line used, with the 3 lines indicated

3
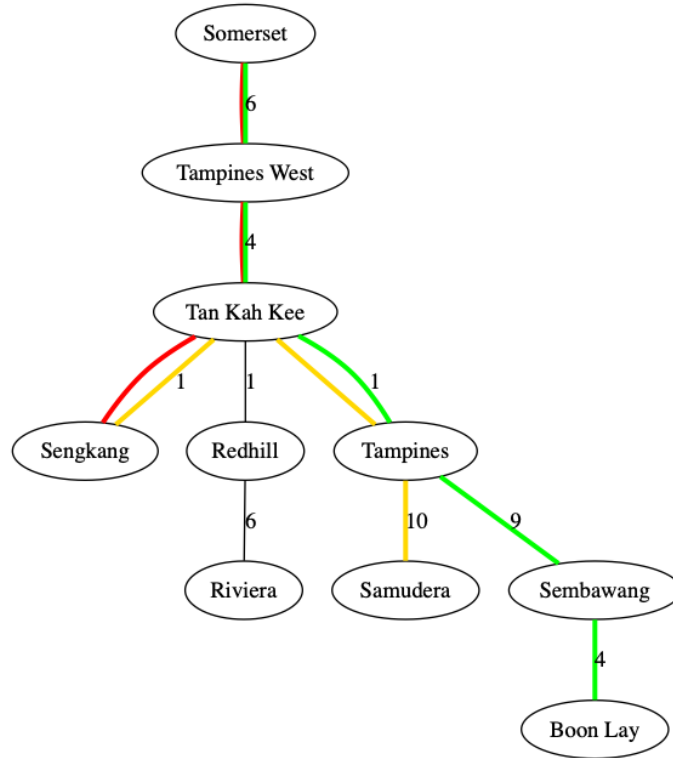
3.2 Graphs

Extreme outliers were removed from the scatter diagram.

Figure 3: Execution Time for 100 ticks

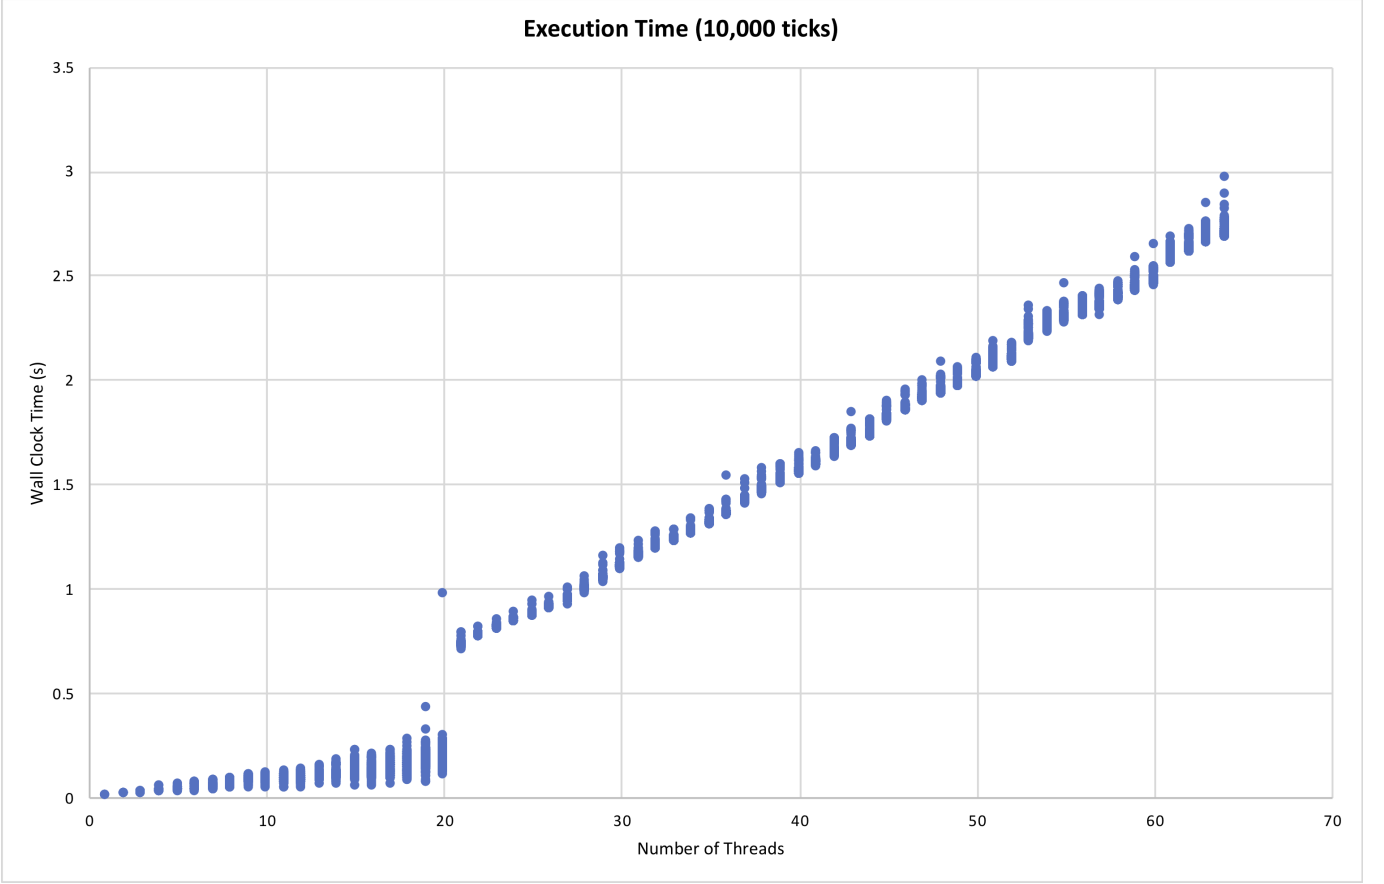Figure 4: Execution Time for 1,000 ticks

Figure 5: Execution Time for 10,000 ticks

Execution Time (10,000 ticks)

4 Discussion

It can be observed that the wall clock time taken for the simulation to complete increases as number of threads increases. This is because as number of threads increases, there is more contention of resources – in our case contention for each train (thread) in using links (edges) between stations (vertices) and contention for each train in opening door at each station. We are managing this using an implicit queue through implementing a timekeeper to track the next allowed event to occur, protected by marking that section as critical. As such, for each link (edge) or station (vertex), only one train (thread) is able to register itself to the timekeeper at any given time – others will have to wait.

However, we have an interesting observation as well. Notice that the execution time behaves very differently when the number of threads is beyond the number of logical cores (20 cores). For small input size (100 ticks), when the number of threads is beyond the number of logical cores, the execution time actually falls. However, as input size gets larger (1,000 and 10,000 ticks), execution time increases. This can be explained that when the number of threads are below the number of logical cores, all the threads are running concurrently, resulting in more contention. However, when number of threads is above the number of logical cores, the threads take turns to wake up and do work, resulting in less contention. For smaller input size, the contention time actually outweighs the execution time, resulting in the fall in execution time. However, for larger input size, there is a large overhead in context-switching which outweighs the effect of contention time. This is supported by the data we collected on number of context-switches for 100 ticks and 1,000 ticks.

We also observe another trend – that the variance in execution time falls when the number of threads exceed the number of logical cores. We currently have no explanation on this, but we suspect that the compiler does an optimisation when the number of threads exceed the number of logical

cores.

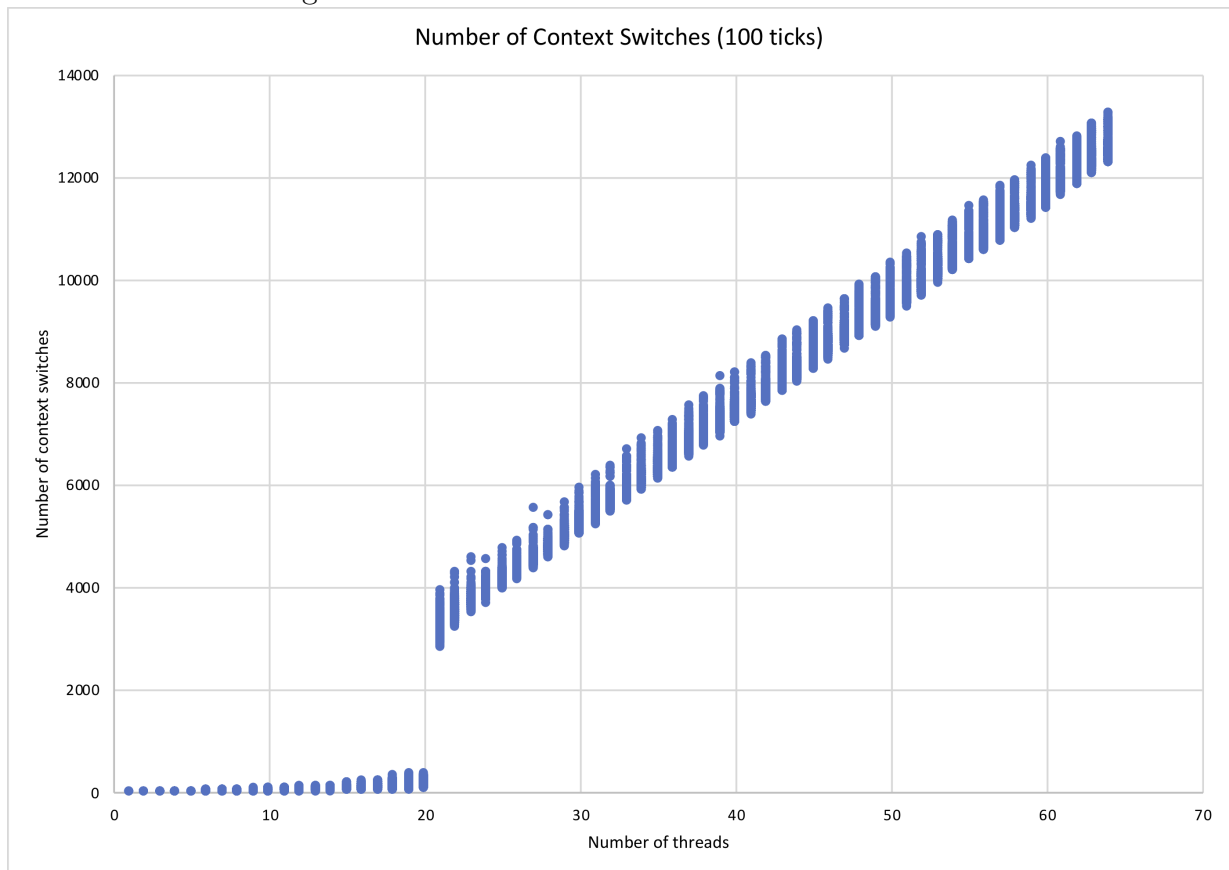Figure 6: Number of context switches for 100 ticks

Number of Context Switches (100 ticks)

Figure 7: Number of context switches for 1,000 ticks

Number of Context Switches (1,000 ticks)

5 Bonus

Starvation will never occur in the simulation program that we wrote. This is because to decide which train to open door or to be allowed to use a link next, we are using an implicit queue to implement First-Come-First-Serve (FCFS) scheduler. As such, every train is assured access to the link or permission to open door aftera long enough time. The assumptions are that no train open its doors or travel using the links indefinitely (which we believe are fair to make).

6 Appendix A: Ruby script used to generate test cases

Below is the code listing of the ruby script used to generate test cases. Essentially, it does:

1. Create a random adjacency matrix with diagonal = 0

2. Find the MST of the random graph created

3. Ensure that there are enough vertices with degree = 1, else go back to step 1

4. Enumerate the 2-combinations of the vertices with degree = 1, and pick 3 randomly.

5. For each of the three 2-combinations, assign them to be the termini of each line.

6. Using breadth-first-search, find the path between the two vertices for each pair of termini.

7. Ensure that the path is long enough, else go back to step 4.

```ruby
require 'set'

MIN_NUM_TERMINI = 4
MIN_NUM_STATIONS_LINE = 4
MRT_STATION_NAMES = ["Jurong East", "Bukit Batok", "Bukit Gombak", "Choa Chu Kang", "Yew
    Tee", "Kranji", "Marsiling", "Woodlands", "Admiralty", "Sembawang", "Yishun",
    "Khatib", "Yio Chu Kang", "Ang Mo Kio", "Bishan", "Braddell", "Toa Payoh", "Novena",
    "Newton", "Orchard", "Somerset", "Dhoby Ghaut", "City Hall", "Raffles Place", "Marina
    Bay", "Marina South Pier", "Pasir Ris", "Tampines", "Simei", "Tanah Merah", "Bedok",
    "Kembangan", "Eunos", "Paya Lebar", "Aljunied", "Kallang", "Lavender", "Bugis", "City
    Hall", "Raffles Place", "Tanjong Pagar", "Outram Park", "Tiong Bahru", "Redhill",
    "Queenstown", "Commonwealth", "Buona Vista", "Dover", "Clementi", "Jurong East",
    "Chinese Garden", "Lakeside", "Boon Lay", "Pioneer", "Joo Koon", "Gul Circle", "Tuas
    Crescent", "Tuas West Road", "Tuas Link", "Expo", "Changi Airport", "HarbourFront",
    "Outram Park", "Chinatown", "Clarke Quay", "Dhoby Ghaut", "Little India", "Farrer
    Park", "Boon Keng", "Potong Pasir", "Woodleigh", "Serangoon", "Kovan", "Hougang",
    "Buangkok", "Sengkang", "Punggol", "Dhoby Ghaut", "Bras Basah", "Esplanade",
    "Promenade", "Nicoll Highway", "Stadium", "Mountbatten", "Dakota", "Paya Lebar",
    "MacPherson", "Tai Seng", "Bartley", "Serangoon", "Lorong Chuan", "Bishan",
    "Marymount", "Caldecott", "Botanic Gardens", "Farrer Road", "Holland Village", "Buona
    Vista", "one-north", "Kent Ridge", "Haw Par Villa", "Pasir Panjang", "Labrador Park",
    "Telok Blangah", "HarbourFront", "Bayfront", "Marina Bay", "Bukit Panjang", "Cashew",
    "Hillview", "Beauty World", "King Albert Park", "Sixth Avenue", "Tan Kah Kee",
    "Botanic Gardens", "Stevens", "Newton", "Little India", "Rochor", "Bugis",
    "Promenade", "Bayfront", "Downtown", "Telok Ayer", "Chinatown", "Fort Canning",
    "Bencoolen", "Jalan Besar", "Bendemeer", "Geylang Bahru", "Mattar", "MacPherson",
    "Ubi", "Kaki Bukit", "Bedok North", "Bedok Reservoir", "Tampines West", "Tampines",
    "Tampines East", "Upper Changi", "Expo", "Choa Chu Kang", "South View", "Keat Hong",
    "Teck Whye", "Phoenix", "Bukit Panjang", "Petir", "Pending", "Bangkit", "Fajar",
    "Segar", "Jelapang", "Senja", "Ten Mile Junction", "Sengkang", "Compassvale",
    "Rumbia", "Bakau", "Kangkar", "Ranggung", "Cheng Lim", "Farmway", "Kupang",
    "Thanggam", "Fernvale", "Layar", "Tongkang", "Renjong", "Punggol", "Cove",
    "Meridian", "Coral Edge", "Riviera", "Kadaloor", "Oasis", "Damai", "Sam Kee", "Teck
    Lee", "Punggol Point", "Samudera", "Nibong", "Sumang", "Soo Teck"]

def generate_random_graph(s, max_weight)
  Array.new(s) { |i| Array.new(s) { |j| i == j ? 0 : rand(1..max_weight) } }
```

```ruby
 9    end
10
11    def print_graph(matrix)
12      matrix.map { |row| row.join(' ') }.join("\n")
13    end
14
15    def prim(matrix)
16      cost = Array.new(matrix.length, Float::INFINITY)
17      parent = Array.new(matrix.length, nil)
18      visited = Array.new(matrix.length, false)
19
20      # start from the first vertex
21      cost[0] = 0
22      parent[0] = -1
23
24      matrix.length.times do
25        u = nil
26        min_weight = Float::INFINITY
27
28        # Find unvisited vertex with minimum cost
29        cost.zip(visited).each_with_index do |zipped, i|
30          c, v = zipped
31          if c < min_weight and !v
32            min_weight = c
33            u = i
34          end
35        end
36        visited[u] = true
37
38        matrix[u].zip(cost, visited).each_with_index do |zipped, i|
39          m, c, v = zipped
40          if m > 0 && !v && c > m
41            cost[i] = m
42            parent[i] = u
43          end
44        end
45      end
46
47      result = Array.new(matrix.length) { Array.new(matrix.length, 0) }
48
49      (1...matrix.length).each do |i|
50        result[i][parent[i]] = result[parent[i]][i] = matrix[i][parent[i]]
51      end
52
53      result
54    end
55
56    def bfs(matrix, termini)
57      from, to = termini
58      open_set = []
59      closed_set = Set[]
60      meta = {}
61
```

```ruby
62      root = from
63      meta[root] = nil
64      open_set.unshift(root)
65
66      while !open_set.empty? do
67        subtree_root = open_set.shift
68        if subtree_root == to
69          return construct_path(subtree_root, meta)
70        end
71        matrix[subtree_root].each_with_index.select { |w, i| w > 0 }.map { |x| x.last }.each
          ↪  do |child|
72          next if closed_set.include?(child)
73          if !open_set.include?(child)
74            meta[child] = subtree_root
75            open_set.unshift(child)
76          end
77        end
78        closed_set.add(subtree_root)
79      end
80    end
81
82    def construct_path(state, meta)
83      result = [state]
84      while !meta[state].nil? do
85        state = meta[state]
86        result.append(state)
87      end
88      result.reverse
89    end
90
91    def permutate_sum(n)
92      (0..n).to_a.flat_map { |i| (0..(n - i)).to_a.map { |j| [i, j, n - i - j] } }
93    end
94
95    def usage_message
96      puts "Invalid args"
97      puts "Usage: ruby test_case_generator.rb <num_vertex> <max_weight> <num_tick>"
98      exit 1
99    end
100
101   # Start of main
102   if ARGV.length != 3
103     usage_message
104   end
105
106   s, max_weight, tick = ARGV.map { |a| a.to_i }
107
108   if s <= 0 || max_weight <= 0
109     usage_message
110   end
111
112   primmed = nil
113   termini = []
```

```ruby
114
115  while termini.length < MIN_NUM_TERMINI do
116    graph = generate_random_graph(s, max_weight)
117    primmed = prim(graph)
118    termini = primmed
119      .each_with_index.select do |row, i|
120        row.reduce(0) { |acc, weight| acc += weight > 0 ? 1 : 0 } == 1
121      end
122      .map { |pair| pair.last }
123  end
124
125  stations = MRT_STATION_NAMES.sample(s)
126  popularities = Array.new(s) { rand(1..10) / 10.0 }
127
128  green_line = []
129  yellow_line = []
130  blue_line = []
131  while green_line.length < MIN_NUM_STATIONS_LINE || yellow_line.length <
    ↪  MIN_NUM_STATIONS_LINE || blue_line.length < MIN_NUM_STATIONS_LINE do
132    green_termini, yellow_termini, blue_termini = termini.combination(2).to_a.sample(3)
133
134    green_line = bfs(primmed, green_termini)
135    yellow_line = bfs(primmed, yellow_termini)
136    blue_line = bfs(primmed, blue_termini)
137  end
138
139  dir_name = "test-#{Time.now.strftime("%Y%m%d-%H%M")}"
140  Dir.mkdir(dir_name)
141
142  (1..64).each do |n|
143    puts "Generating test cases for #{n} threads"
144    permutate_sum(n).each do |trains|
145      File.open("#{dir_name}/#{trains.join("-")}", "w") do |f|
146        f.puts s
147        f.puts stations.join(",")
148        f.puts print_graph(primmed)
149        f.puts popularities.join(",")
150        f.puts green_line.map { |s| stations[s] }.join(",")
151        f.puts yellow_line.map { |s| stations[s] }.join(",")
152        f.puts blue_line.map { |s| stations[s] }.join(",")
153        f.puts tick
154        f.puts trains.join(",")
155      end
156    end
157  end
```